

major-code-pdf

May 27, 2024

```
[1]: import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt
sb.set()
from matplotlib import pyplot as plt, font_manager as fm
```

```
[2]: milkData = pd.read_csv("C:\\Users\\harsh\\Downloads\\milknew (2).csv")
milkData
```

```
[2]:
```

	pH	Temprature	Taste	Odor	Fat	Turbidity	Colour	Grade
0	6.6	35	1	0	1	0	254	high
1	6.6	36	0	1	0	1	253	high
2	8.5	70	1	1	1	1	246	low
3	9.5	34	1	1	0	1	255	low
4	6.6	37	0	0	0	0	255	medium
...
1054	6.7	45	1	1	0	0	247	medium
1055	6.7	38	1	0	1	0	255	high
1056	3.0	40	1	1	1	1	255	low
1057	6.8	43	1	0	1	0	250	high
1058	8.6	55	0	1	1	1	255	low

[1059 rows x 8 columns]

```
[3]: milkData.head(5)
```

```
[3]:
```

	pH	Temprature	Taste	Odor	Fat	Turbidity	Colour	Grade
0	6.6	35	1	0	1	0	254	high
1	6.6	36	0	1	0	1	253	high
2	8.5	70	1	1	1	1	246	low
3	9.5	34	1	1	0	1	255	low
4	6.6	37	0	0	0	0	255	medium

```
[4]: milkData.tail(5)
```

```
[4]:
```

	pH	Temprature	Taste	Odor	Fat	Turbidity	Colour	Grade
1054	6.7	45	1	1	0	0	247	medium
1055	6.7	38	1	0	1	0	255	high
1056	3.0	40	1	1	1	1	255	low
1057	6.8	43	1	0	1	0	250	high
1058	8.6	55	0	1	1	1	255	low

```
[5]: milkData.keys()
```

```
[5]: Index(['pH', 'Temprature', 'Taste', 'Odor', 'Fat ', 'Turbidity', 'Colour',
          'Grade'],
          dtype='object')
```

```
[6]: milkData.describe()
```

```
[6]:
```

	pH	Temprature	Taste	Odor	Fat \
count	1059.000000	1059.000000	1059.000000	1059.000000	1059.000000
mean	6.630123	44.226629	0.546742	0.432483	0.671388
std	1.399679	10.098364	0.498046	0.495655	0.469930
min	3.000000	34.000000	0.000000	0.000000	0.000000
25%	6.500000	38.000000	0.000000	0.000000	0.000000
50%	6.700000	41.000000	1.000000	0.000000	1.000000
75%	6.800000	45.000000	1.000000	1.000000	1.000000
max	9.500000	90.000000	1.000000	1.000000	1.000000

	Turbidity	Colour
count	1059.000000	1059.000000
mean	0.491029	251.840415
std	0.500156	4.307424
min	0.000000	240.000000
25%	0.000000	250.000000
50%	0.000000	255.000000
75%	1.000000	255.000000
max	1.000000	255.000000

```
[7]: milkData.shape
```

```
[7]: (1059, 8)
```

```
[8]: milkData.dtypes
```

```
[8]: pH                float64
     Temprature        int64
     Taste             int64
     Odor              int64
     Fat               int64
     Turbidity          int64
```

```
Colour          int64
Grade           object
dtype: object
```

```
[9]: milkData.dropna()
milkData.shape
```

```
[9]: (1059, 8)
```

```
[10]: milkData.duplicated().sum()
```

```
[10]: 976
```

```
[11]: milkData.loc[milkData.duplicated(),:]
```

```
[11]:
```

	pH	Temprature	Taste	Odor	Fat	Turbidity	Colour	Grade
35	6.8	45	0	1	1	1	255	high
48	9.5	34	1	1	0	1	255	low
50	6.6	37	1	1	1	1	255	high
51	5.5	45	1	0	1	1	250	low
52	4.5	60	0	1	1	1	250	low
...
1054	6.7	45	1	1	0	0	247	medium
1055	6.7	38	1	0	1	0	255	high
1056	3.0	40	1	1	1	1	255	low
1057	6.8	43	1	0	1	0	250	high
1058	8.6	55	0	1	1	1	255	low

```
[976 rows x 8 columns]
```

```
[12]: milkData.isnull().sum()
```

```
[12]: pH          0
Temprature      0
Taste          0
Odor           0
Fat            0
Turbidity      0
Colour         0
Grade         0
dtype: int64
```

```
[13]: milkData.nunique()
```

```
[13]: pH          16
Temprature      17
Taste          2
```

```

Odor          2
Fat           2
Turbidity     2
Colour        9
Grade         3
dtype: int64

```

```

[14]: for i in milkData.columns:
        print(i)
        print(milkData[i].unique())
        print('-----')

```

```

pH
[6.6 8.5 9.5 5.5 4.5 8.1 6.7 5.6 8.6 7.4 6.8 6.5 4.7 3.  9.  6.4]
-----
Temprature
[35 36 70 34 37 45 60 66 50 55 90 38 40 43 42 41 65]
-----
Taste
[1 0]
-----
Odor
[0 1]
-----
Fat
[1 0]
-----
Turbidity
[0 1]
-----
Colour
[254 253 246 255 250 247 245 240 248]
-----
Grade
['high' 'low' 'medium']
-----

```

```

[15]: for i in milkData.columns:
        print(i)
        print(milkData[i].value_counts())
        print('-----')

```

```

pH
pH
6.8    249
6.5    189
6.6    159
6.7     82

```

3.0	70
9.0	61
8.6	40
7.4	39
4.5	37
9.5	24
8.1	24
5.5	23
8.5	22
4.7	20
5.6	19
6.4	1

Name: count, dtype: int64

Temprature
Temprature

45	219
38	179
40	132
37	83
43	77
36	66
50	58
55	48
34	40
41	30
66	24
35	23
70	22
65	22
60	18
90	17
42	1

Name: count, dtype: int64

Taste
Taste

1	579
0	480

Name: count, dtype: int64

Odor
Odor

0	601
1	458

Name: count, dtype: int64

Fat

```
Fat
1    711
0    348
Name: count, dtype: int64
```

```
Turbidity
Turbidity
0    539
1    520
Name: count, dtype: int64
```

```
Colour
Colour
255    628
250    146
245    115
247     48
246     44
240     32
248     23
253     22
254      1
Name: count, dtype: int64
```

```
Grade
Grade
low      429
medium   374
high     256
Name: count, dtype: int64
```

```
[16]: milkData.info
```

```
[16]: <bound method DataFrame.info of          pH  Temprature  Taste  Odor  Fat
Turbidity  Colour    Grade
0      6.6          35     1     0     1         0      254   high
1      6.6          36     0     1     0         1      253   high
2      8.5          70     1     1     1         1      246   low
3      9.5          34     1     1     0         1      255   low
4      6.6          37     0     0     0         0      255  medium
...  ...          ...  ...  ...  ...  ...  ...
1054  6.7          45     1     1     0         0      247  medium
1055  6.7          38     1     0     1         0      255   high
1056  3.0          40     1     1     1         1      255   low
1057  6.8          43     1     0     1         0      250   high
1058  8.6          55     0     1     1         1      255   low
```

```
[1059 rows x 8 columns]>
```

```
[17]: milkData1= pd.DataFrame(milkData)
```

```
[18]: milkData2=milkData.drop(['Grade'], axis=1)
```

```
[19]: milkData2
```

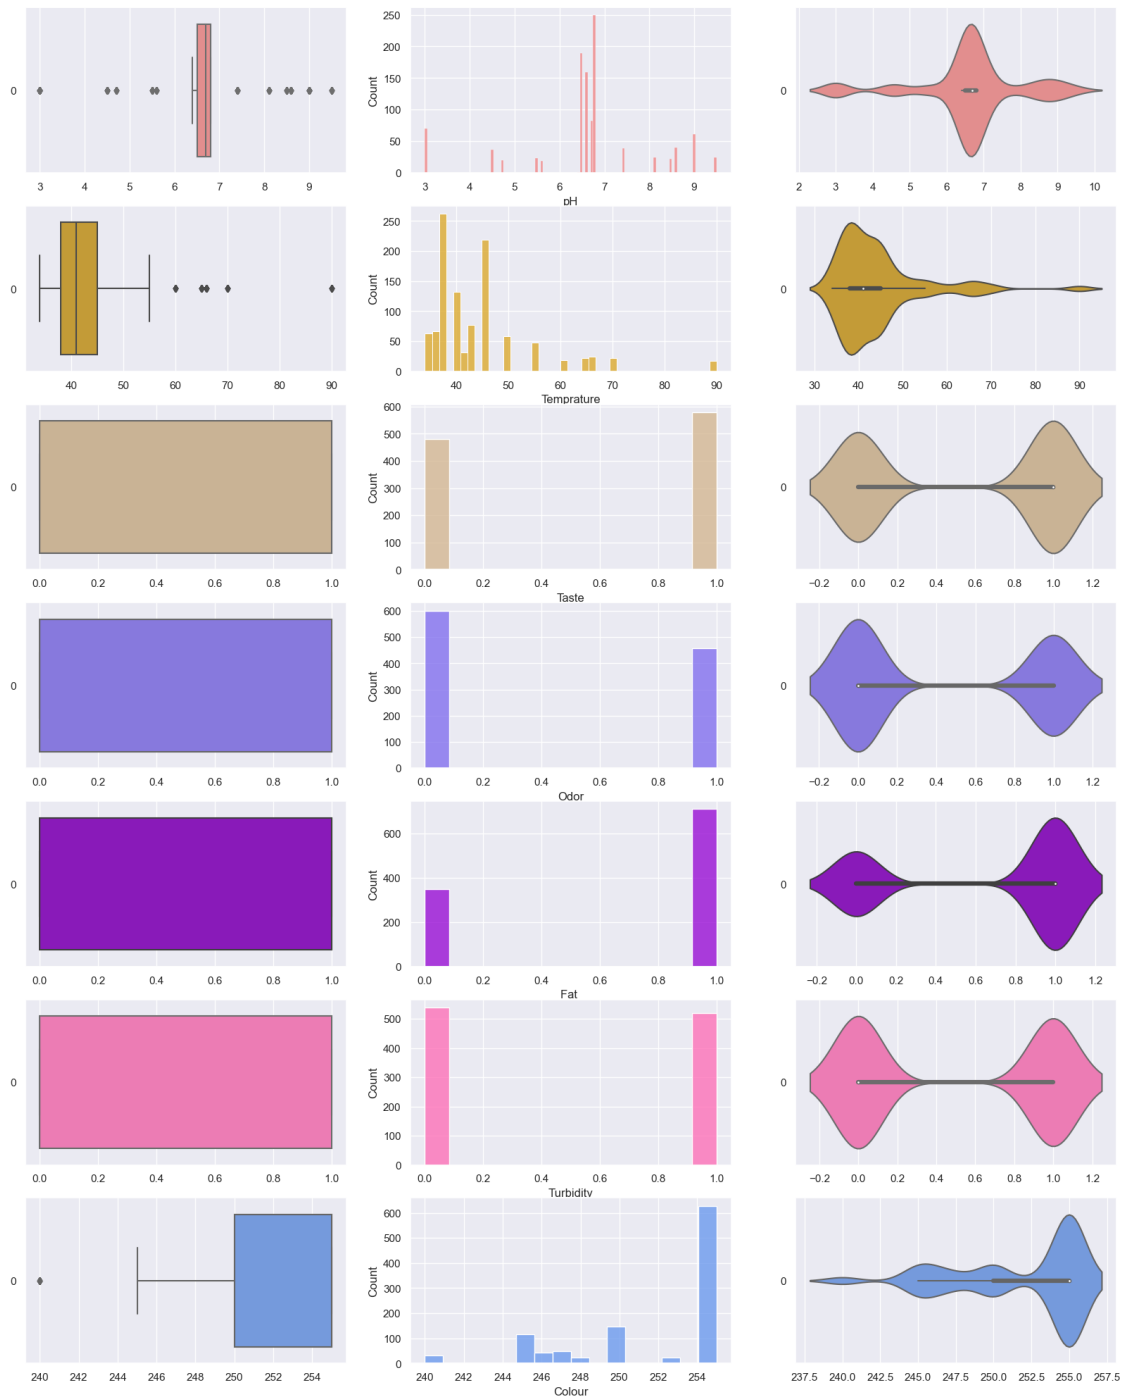
```
[19]:
```

	pH	Temprature	Taste	Odor	Fat	Turbidity	Colour
0	6.6	35	1	0	1	0	254
1	6.6	36	0	1	0	1	253
2	8.5	70	1	1	1	1	246
3	9.5	34	1	1	0	1	255
4	6.6	37	0	0	0	0	255
...
1054	6.7	45	1	1	0	0	247
1055	6.7	38	1	0	1	0	255
1056	3.0	40	1	1	1	1	255
1057	6.8	43	1	0	1	0	250
1058	8.6	55	0	1	1	1	255

```
[1059 rows x 7 columns]
```

```
[20]: f, axes = plt.subplots(7, 3, figsize=(20, 25))
colors = ["lightcoral", "goldenrod", "tan", "mediumslateblue", "darkviolet",
↪ "hotpink", "cornflowerblue", "dodgerblue", "royalblue", "mediumaquamarine",
↪ "teal", "firebrick"]

count = 0
for var in milkData2:
    sb.boxplot(data=milkData2[var], orient = "h", color = colors[count], ax =
↪ axes[count,0])
    sb.histplot(data=milkData2[var], color = colors[count], ax = axes[count,1])
    sb.violinplot(data=milkData2[var], orient = "h", color = colors[count], ax=
↪ axes[count,2])
    count += 1
```

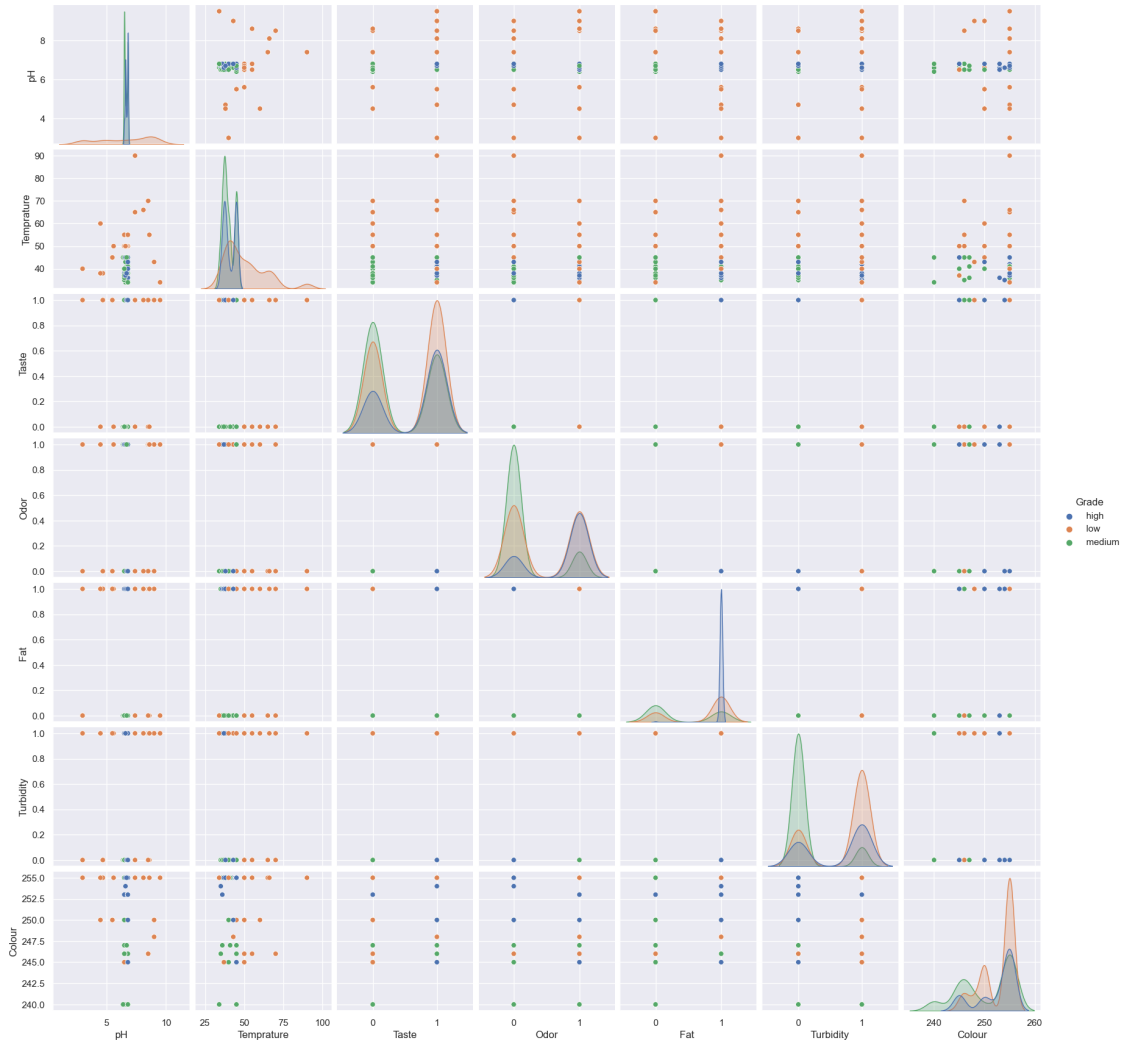


```
[21]: sb.pairplot(milkData,vars=milkData.columns[:-1],hue='Grade')
```


C:\Users\harsh\OneDrive\Documents\anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

```
self._figure.tight_layout(*args, **kwargs)
```


[21]: <seaborn.axisgrid.PairGrid at 0x24904625f50>



```
[22]: cols = milkData.columns[:6]
densityplot = milkData[cols].plot(kind='density')
```



```
[23]: print('The lowest temperature in the Milk is ' + str(milkData['Temprature'].
        ↳min()))
print('The highest temperature in the Milk is ' + str(milkData['Temprature'].
        ↳max()))
print('The average temperature in the Milk is ' +
        ↳str(round(milkData['Temprature'].mean(),3)))
#min, max, average for PH
print('The lowest pH in the Milk is ' + str(milkData['pH'].min()))
print('The highest pH in the Milk is ' + str(milkData['pH'].max()))
print('The average pH in the Milk is ' + str(round(milkData['pH'].mean(),3)))
#min,max,averagr for colour
print('The lowest Colour in the Milk is ' + str(milkData['Colour'].min()))
print('The highest Colour in the Milk is ' + str(milkData['Colour'].max()))
print('The average Colour in the Milk is ' + str(round(milkData['Colour'].
        ↳mean(),3)))
```

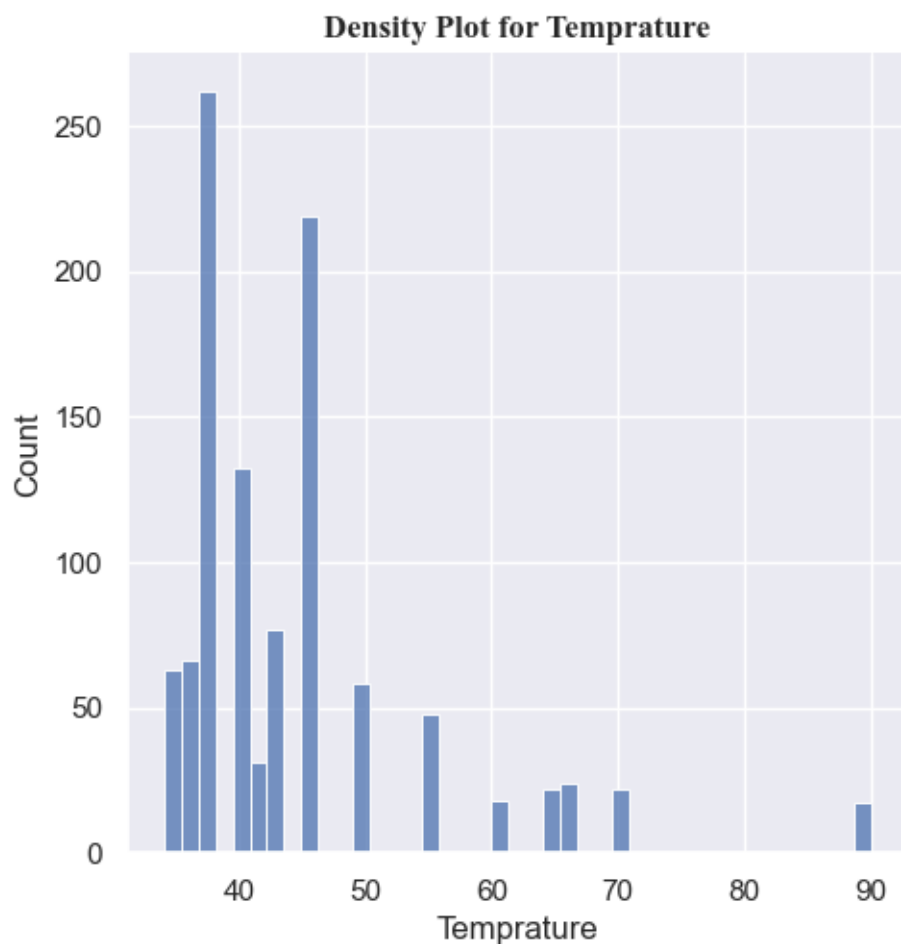
```
The lowest temperature in the Milk is 34
The highest temperature in the Milk is 90
The average temperature in the Milk is 44.227
The lowest pH in the Milk is 3.0
The highest pH in the Milk is 9.5
The average pH in the Milk is 6.63
```

The lowest Colour in the Milk is 240
The highest Colour in the Milk is 255
The average Colour in the Milk is 251.84

```
[24]: sb.displot(milkData['Temprature'], color = 'b')
plt.title("Density Plot for Temprature",fontname="Times New Roman", size=12,
↵,fontweight="bold" )
plt.show()
```

C:\Users\harsh\OneDrive\Documents\anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

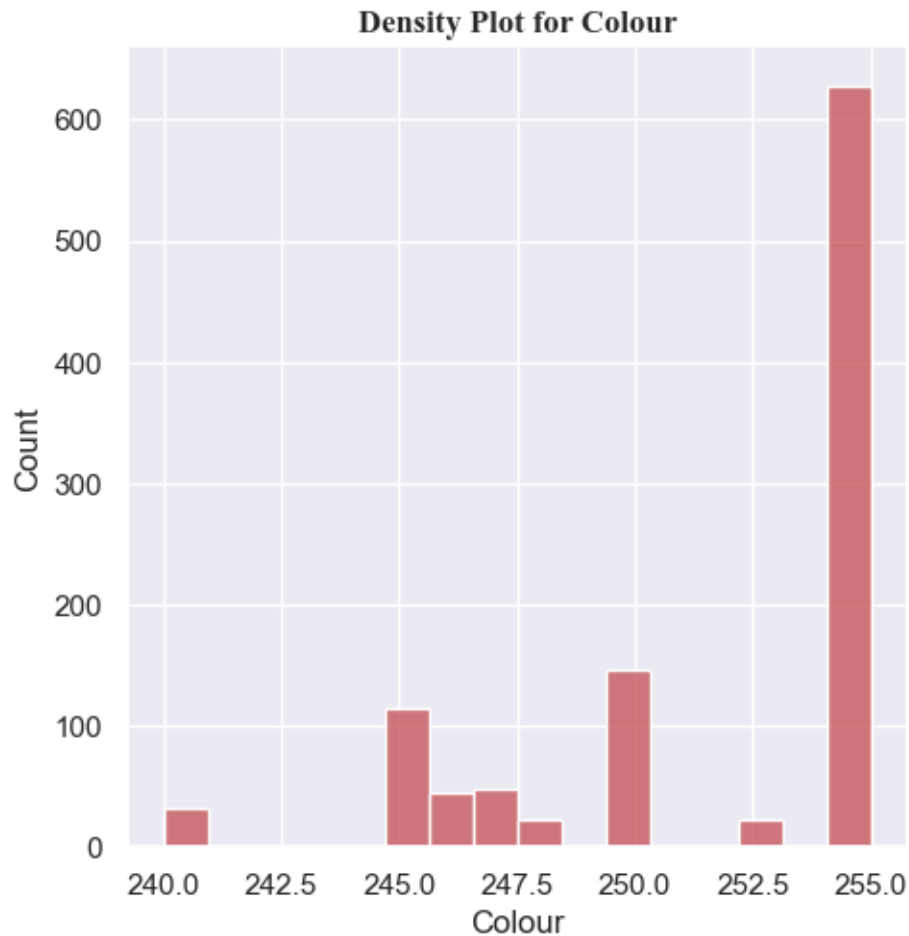
```
self._figure.tight_layout(*args, **kwargs)
```



```
[25]: sb.displot(milkData['Colour'], color = 'r')
plt.title("Density Plot for Colour",fontname="Times New Roman", size=12,
↵,fontweight="bold" )
plt.show()
```

C:\Users\harsh\OneDrive\Documents\anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight


```
self._figure.tight_layout(*args, **kwargs)
```



```
[26]: sb.displot(milkData['pH'], color = 'y')
plt.title("Density Plot for pH",fontname="Times New Roman", size=12,
↪,fontweight="bold" )
plt.show()
```


C:\Users\harsh\OneDrive\Documents\anaconda\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight

```
self._figure.tight_layout(*args, **kwargs)
```



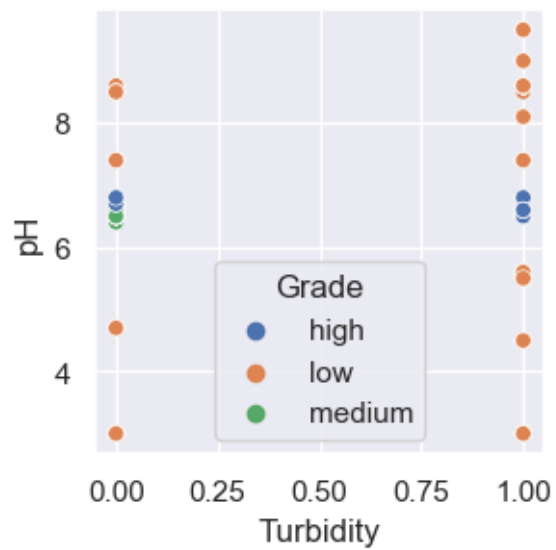
```
[27]: milkData['Taste'].value_counts().plot(kind='pie', autopct='%.2f')
```

```
[27]: <Axes: ylabel='count'>
```



```
[28]: import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(3,3))
sns.scatterplot(x='Turbidity', y='pH', hue='Grade', data=milkData)
plt.show()
```



```
[29]: milkData.Grade.value_counts()
```

```
[29]: Grade  
low      429  
medium   374  
high     256  
Name: count, dtype: int64
```

```
[30]: sb.countplot(x='Grade', data=milkData)  
plt.show()
```



```
[31]: # Exclude non-numeric columns  
numeric_data = milkData.select_dtypes(include=['float64', 'int64'])  
  
# Display the correlation matrix  
correlation_matrix = numeric_data.corr()  
print(correlation_matrix)
```

```


# Plot the heatmap
fig, axes = plt.subplots(1, 1, figsize=(10, 10))
heatmap = sb.heatmap(correlation_matrix, vmin=-1, vmax=1, linewidths=1,
                    annot=True, fmt=".2f", annot_kws={"size": 14}, cmap="bwr")

# Show the plot
plt.show()

```

	pH	Temprature	Taste	Odor	Fat	Turbidity	\
pH	1.000000	0.244684	-0.064053	-0.081331	-0.093429	0.048384	
Temprature	0.244684	1.000000	-0.109792	-0.048870	0.024073	0.185106	
Taste	-0.064053	-0.109792	1.000000	0.017582	0.324149	0.055755	
Odor	-0.081331	-0.048870	0.017582	1.000000	0.314505	0.457935	
Fat	-0.093429	0.024073	0.324149	0.314505	1.000000	0.329264	
Turbidity	0.048384	0.185106	0.055755	0.457935	0.329264	1.000000	
Colour	-0.164565	-0.008511	-0.082654	-0.039361	0.114151	0.136436	

	Colour
pH	-0.164565
Temprature	-0.008511
Taste	-0.082654
Odor	-0.039361
Fat	0.114151
Turbidity	0.136436
Colour	1.000000



```
[32]: pd.DataFrame(milkData.groupby("Odor")["Turbidity"].mean())
```

```
[32]:      Turbidity
Odor
0      0.291181
1      0.753275
```

```
[33]: pd.DataFrame(milkData.groupby('Fat')['Taste'].mean())
```

```
[33]:      Taste
Fat
0      0.316092
```

1 0.659634

```
[34]: pd.DataFrame(milkData.groupby('pH')['Temprature'].mean())
```

```
[34]:      Temprature
pH
3.0    40.000000
4.5    48.702703
4.7    38.000000
5.5    45.000000
5.6    50.000000
6.4    45.000000
6.5    37.746032
6.6    41.125786
6.7    42.951220
6.8    42.273092
7.4    75.897436
8.1    66.000000
8.5    70.000000
8.6    55.000000
9.0    43.000000
9.5    34.000000
```

```
[35]: #transform Grade in numerical numbers using label encoder
from sklearn import preprocessing
label_encoder = preprocessing.LabelEncoder()
milkData['Grade'] = label_encoder.fit_transform(milkData['Grade'])
```

```
[36]: milkData.head()
```

```
[36]:      pH  Temprature  Taste  Odor  Fat  Turbidity  Colour  Grade
0  6.6         35      1    0    1         0     254      0
1  6.6         36      0    1    0         1     253      0
2  8.5         70      1    1    1         1     246      1
3  9.5         34      1    1    0         1     255      1
4  6.6         37      0    0    0         0     255      2
```

```
[37]: from sklearn.linear_model import LogisticRegression #logistic regression
from sklearn.ensemble import RandomForestClassifier #Random Forest
from sklearn.neighbors import KNeighborsClassifier #KNN
from sklearn.tree import DecisionTreeClassifier #Decision Tree
from sklearn.model_selection import train_test_split #training and testing data
↳ split
from sklearn import metrics #accuracy measure
from sklearn.metrics import confusion_matrix #for confusion matrix
```

```
[38]: #splitting into train and test
x= milkData.drop(['Grade'],axis=1)
y= milkData['Grade']

from sklearn.preprocessing import StandardScaler
PredictorScaler=StandardScaler()

PredictorScalerFit=PredictorScaler.fit(x)

x=PredictorScalerFit.transform(x)
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.
↪3,random_state=42)
```

```
[39]: print("X Train : ", X_train.shape)
print("X Test  : ", X_test.shape)
print("Y Train : ", y_train.shape)
print("Y Test  : ", y_test.shape)
```

```
X Train : (741, 7)
X Test  : (318, 7)
Y Train : (741,)
Y Test  : (318,)
```

```
[40]: # Import necessary libraries
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression

# Create a StandardScaler instance
scaler = StandardScaler()

# Fit the scaler to the training data
scaler.fit(X_train)

# Transform the training and test sets
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train a logistic regression model on the scaled data
model = LogisticRegression()
model.fit(X_train_scaled, y_train)
```

```
[40]: LogisticRegression()
```

```
[41]: #This is for the test dataset
model = LogisticRegression()
model.fit(X_train_scaled,y_train)
X_test_predict_lr =model.predict(X_test)
```

```
X_test_predict_lr
```

```
[41]: array([1, 1, 1, 1, 2, 1, 1, 0, 0, 0, 1, 2, 0, 2, 1, 1, 1, 2, 2, 2, 0, 2,
          0, 2, 2, 0, 2, 0, 1, 1, 2, 2, 0, 0, 0, 2, 2, 0, 2, 2, 1, 1, 0, 0,
          0, 0, 0, 0, 2, 0, 2, 1, 0, 2, 2, 1, 0, 1, 1, 2, 2, 1, 1, 0, 0, 0,
          2, 1, 1, 2, 1, 1, 0, 2, 1, 1, 1, 2, 1, 1, 2, 2, 0, 0, 2, 2, 0, 1,
          2, 0, 1, 1, 0, 0, 1, 1, 1, 2, 1, 0, 1, 1, 1, 1, 2, 2, 1, 1, 2, 0,
          1, 2, 2, 1, 0, 1, 2, 1, 2, 0, 2, 1, 1, 1, 2, 1, 2, 2, 0, 2, 1, 2,
          1, 2, 1, 1, 2, 1, 2, 2, 2, 0, 1, 2, 1, 1, 2, 1, 1, 2, 2, 2, 0, 1,
          2, 0, 1, 0, 1, 0, 2, 2, 2, 1, 2, 0, 1, 1, 2, 1, 0, 0, 1, 1, 0, 1,
          2, 2, 2, 2, 1, 2, 0, 1, 0, 1, 2, 0, 2, 1, 1, 2, 0, 1, 0, 1, 1, 1,
          2, 2, 0, 2, 0, 0, 2, 1, 2, 0, 2, 2, 0, 0, 2, 2, 0, 1, 1, 0, 2, 1,
          2, 2, 0, 0, 1, 0, 1, 0, 2, 0, 0, 1, 2, 0, 0, 1, 2, 2, 2, 1, 1, 0,
          1, 1, 1, 0, 1, 2, 1, 1, 1, 0, 2, 0, 2, 2, 1, 2, 1, 0, 2, 0, 0, 0,
          2, 1, 0, 0, 2, 0, 2, 2, 1, 0, 2, 1, 1, 1, 2, 0, 2, 0, 1, 0, 1, 1,
          1, 2, 0, 2, 1, 1, 2, 0, 2, 1, 0, 0, 2, 1, 2, 0, 1, 2, 1, 1, 1, 0,
          2, 2, 2, 0, 2, 2, 2, 1, 1, 0])
```

```
[42]: #This is for the test dataset
model = LogisticRegression()
model.fit(X_train_scaled,y_train)
X_test_predict_lr =model.predict(X_test)
print('The accuracy for Logistic Regression model is (Test Dataset) ', metrics.
      ↪accuracy_score(X_test_predict_lr,y_test))
#This is for the train dataset
model = LogisticRegression()
model.fit(X_test_scaled,y_test)
X_train_predict_lr =model.predict(X_train)
print('The accuracy for Logistic Regression model is (Train Dataset) ',metrics.
      ↪accuracy_score(X_train_predict_lr,y_train))
```

The accuracy for Logistic Regression model is (Test Dataset) 0.8522012578616353
The accuracy for Logistic Regression model is (Train Dataset)
0.8259109311740891

```
[43]: from sklearn.metrics import classification_report

# Assuming 'model_lr' is the Logistic Regression model and 'model_rfc' is the
      ↪Random Forest model
# (Ensure that these models have been trained before running this code)

# Split the data into training and testing sets
x = milkData.drop(['Grade'], axis=1)
y = milkData['Grade']

# Standardize the features
PredictorScaler = StandardScaler()
```

```

PredictorScalerFit = PredictorScaler.fit(x)
x = PredictorScalerFit.transform(x)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
    ↪random_state=42)

# Train the Logistic Regression model
model_lr = LogisticRegression()
model_lr.fit(X_train, y_train)

# Predictions on the testing set
y_pred_lr = model_lr.predict(X_test)

# Logistic Regression classification report
print("Classification Report for Logistic Regression (Test Dataset):")
print(classification_report(y_test, y_pred_lr))

```

Classification Report for Logistic Regression (Test Dataset):

	precision	recall	f1-score	support
0	0.72	0.84	0.78	76
1	0.86	0.87	0.87	115
2	0.94	0.83	0.88	127
accuracy			0.85	318
macro avg	0.84	0.85	0.84	318
weighted avg	0.86	0.85	0.85	318

```

[44]: from sklearn.metrics import confusion_matrix, classification_report,
    ↪ConfusionMatrixDisplay
import matplotlib.pyplot as plt

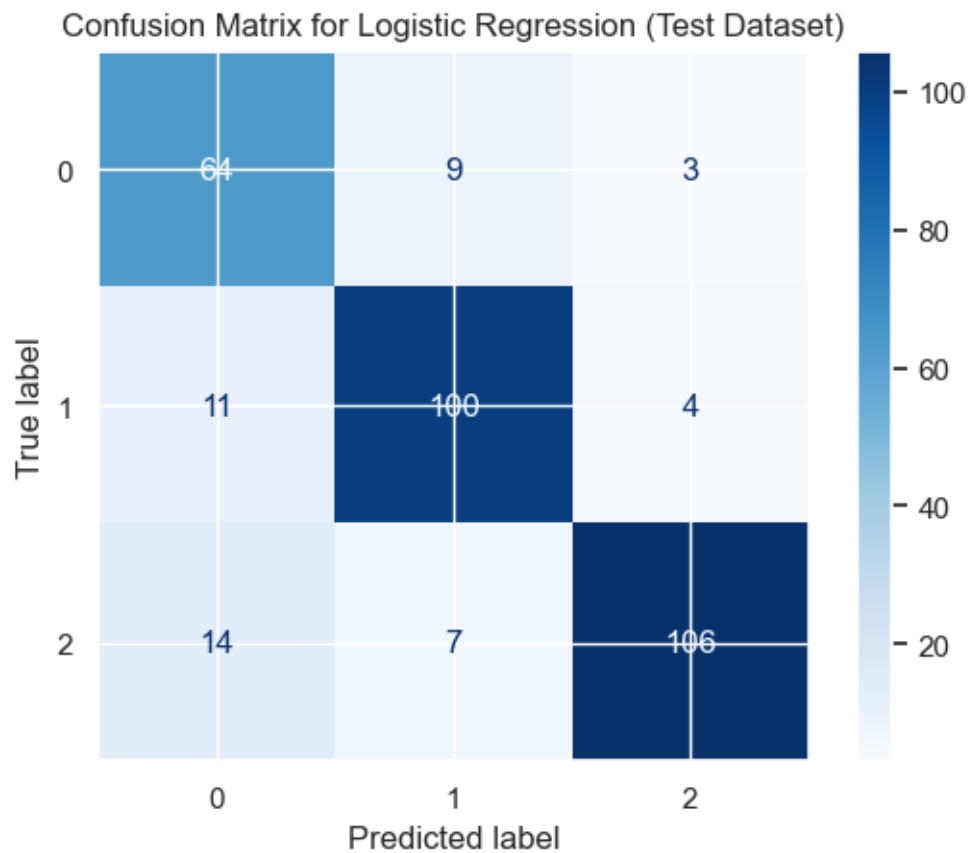
# Assuming 'y_test' and 'y_pred_lr' are the true labels and predicted labels,
    ↪respectively

# Create confusion matrix
conf_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Display confusion matrix using ConfusionMatrixDisplay
disp_lr = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_lr,
    ↪display_labels=model_lr.classes_)
disp_lr.plot(cmap='Blues', values_format='d')

```

```
plt.title('Confusion Matrix for Logistic Regression (Test Dataset)')
plt.show()
```



```
[45]: from sklearn.metrics import classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Assuming 'model_lr' is the Logistic Regression model
# Ensure that 'model_lr' has been trained before running this code

# Split the data into training and testing sets
x = milkData.drop(['Grade'], axis=1)
y = milkData['Grade']

# Standardize the features
PredictorScaler = StandardScaler()
PredictorScalerFit = PredictorScaler.fit(x)
x = PredictorScalerFit.transform(x)
```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
↳random_state=42)

# Train the Logistic Regression model
model_lr = LogisticRegression()
model_lr.fit(X_train, y_train)

# Predictions on the training set
y_pred_lr_train = model_lr.predict(X_train)

# Logistic Regression classification report for training dataset
print("Classification Report for Logistic Regression (Training Dataset):")
print(classification_report(y_train, y_pred_lr_train))

```

Classification Report for Logistic Regression (Training Dataset):

	precision	recall	f1-score	support
0	0.72	0.84	0.78	180
1	0.88	0.87	0.87	314
2	0.93	0.84	0.89	247
accuracy			0.85	741
macro avg	0.85	0.85	0.85	741
weighted avg	0.86	0.85	0.85	741


```

[46]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Compute the confusion matrix
conf_matrix_train = confusion_matrix(y_train, y_pred_lr_train)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues',
            xticklabels=model_lr.classes_, yticklabels=model_lr.classes_)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Logistic Regression (Training Dataset)')
plt.show()

```



```
[47]: model = RandomForestClassifier(n_estimators=100)
model.fit(X_train, y_train)
# Make predictions on the test dataset
X_test_predict_rfc = model.predict(X_test)
X_test_predict_rfc
```

```
[47]: array([0, 1, 1, 2, 2, 1, 1, 0, 1, 2, 1, 2, 0, 2, 1, 1, 0, 2, 2, 2, 1, 2,
0, 2, 2, 2, 0, 0, 1, 1, 2, 0, 0, 2, 0, 2, 2, 0, 2, 2, 1, 1, 1, 2,
0, 1, 0, 0, 2, 0, 2, 1, 1, 2, 2, 1, 0, 1, 1, 2, 2, 1, 0, 0, 0, 1,
2, 2, 1, 2, 1, 1, 0, 2, 1, 1, 1, 2, 1, 2, 2, 2, 0, 0, 2, 2, 2, 1,
2, 0, 1, 1, 0, 2, 1, 2, 0, 0, 1, 0, 1, 1, 1, 1, 2, 2, 1, 1, 2, 0,
0, 2, 2, 1, 0, 1, 2, 1, 1, 0, 2, 1, 1, 1, 2, 0, 2, 2, 0, 2, 1, 2,
1, 2, 1, 1, 2, 1, 2, 2, 2, 0, 1, 2, 1, 1, 0, 1, 1, 2, 2, 2, 2, 1,
2, 1, 1, 0, 1, 0, 2, 2, 2, 1, 2, 0, 1, 1, 2, 1, 2, 0, 1, 1, 0, 1,
2, 2, 2, 2, 1, 2, 0, 1, 0, 1, 2, 0, 2, 1, 1, 2, 2, 1, 0, 1, 1, 0,
2, 2, 2, 2, 0, 0, 2, 1, 2, 2, 2, 2, 0, 0, 2, 2, 0, 1, 1, 0, 2, 1,
2, 2, 0, 0, 1, 0, 1, 1, 2, 0, 0, 1, 2, 1, 2, 1, 2, 1, 1, 0,
0, 1, 1, 0, 1, 2, 1, 1, 1, 0, 2, 0, 2, 2, 1, 2, 1, 1, 2, 0, 0, 0,
```



```

2, 1, 0, 0, 2, 2, 2, 2, 1, 0, 2, 1, 1, 1, 2, 0, 2, 0, 2, 0, 1, 1,
1, 2, 0, 2, 1, 1, 2, 0, 2, 2, 0, 0, 2, 1, 2, 1, 0, 2, 1, 1, 1, 0,
2, 2, 2, 2, 2, 2, 2, 2, 2, 1, 0])

```

```

[48]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import accuracy_score
      model = RandomForestClassifier(n_estimators=100)
      model.fit(X_train, y_train)
      X_test_predict_rfc = model.predict(X_test)
      print('The accuracy of the Random Forests model is (Test Dataset)',
            ↪accuracy_score(X_test_predict_rfc, y_test))
      model_train = RandomForestClassifier(n_estimators=100)
      model_train.fit(X_test, y_test)
      X_train_predict_rfc = model_train.predict(X_train)
      print('The accuracy of the Random Forests model is (Train Dataset)',
            ↪accuracy_score(X_train_predict_rfc, y_train))

```

The accuracy of the Random Forests model is (Test Dataset) 0.9968553459119497
The accuracy of the Random Forests model is (Train Dataset) 0.9986504723346828

```

[49]: model = RandomForestClassifier(n_estimators=100)
      model.fit(X_train, y_train)

      # Predict on the test set
      y_test_predict_rfc = model.predict(X_test)

      # Create confusion matrix
      conf_matrix = confusion_matrix(y_test, y_test_predict_rfc)

      # Print Classification Report
      print('Classification Report for Random Forest (Test Dataset):\n',
            ↪classification_report(y_test, y_test_predict_rfc))

```

Classification Report for Random Forest (Test Dataset):

	precision	recall	f1-score	support
0	0.99	1.00	0.99	76
1	1.00	0.99	1.00	115
2	1.00	1.00	1.00	127
accuracy			1.00	318
macro avg	1.00	1.00	1.00	318
weighted avg	1.00	1.00	1.00	318

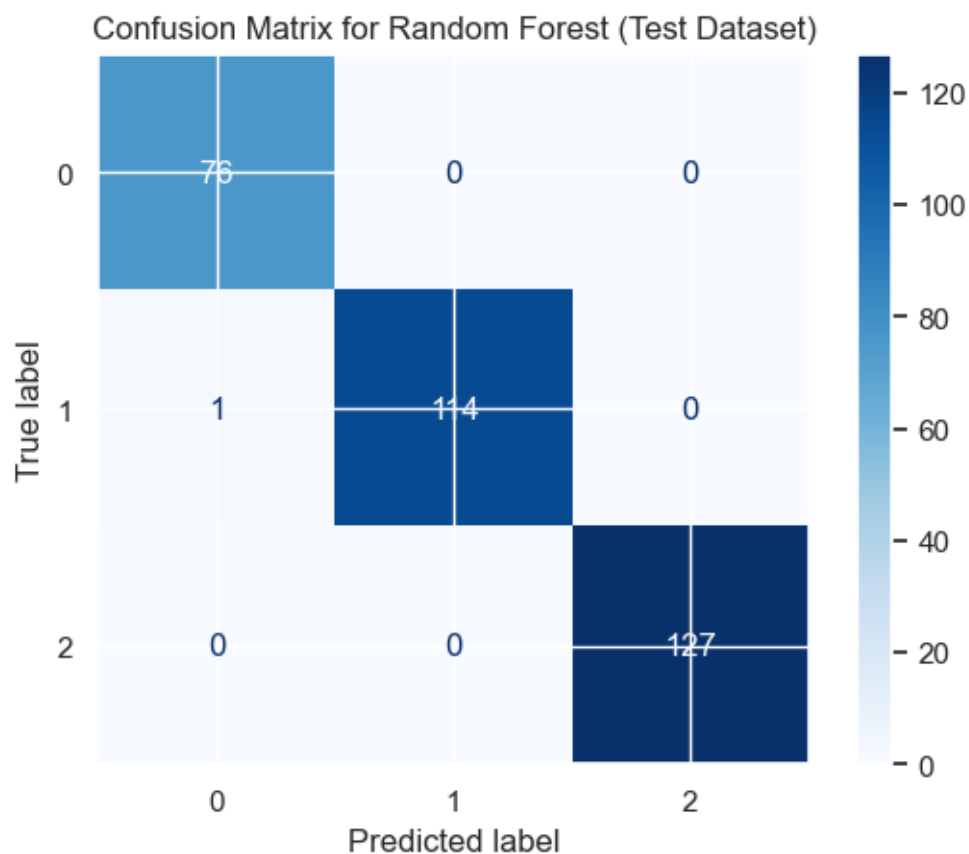
```
[50]: from sklearn.metrics import confusion_matrix, classification_report,
      ↪ ConfusionMatrixDisplay
      import matplotlib.pyplot as plt

      # Assuming 'y_test' and 'y_test_predict_rfc' are the true labels and predicted
      ↪ labels, respectively

      # Create confusion matrix
      conf_matrix = confusion_matrix(y_test, y_test_predict_rfc)

      # Display confusion matrix using ConfusionMatrixDisplay
      disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix,
      ↪ display_labels=model.classes_)
      disp.plot(cmap='Blues', values_format='d')

      plt.title('Confusion Matrix for Random Forest (Test Dataset)')
      plt.show()
```



```
[51]: from sklearn.metrics import confusion_matrix, classification_report
import numpy as np

# Predict on the training set
y_train_predict_rfc = model.predict(X_train)

# Create confusion matrix for training dataset
conf_matrix_train = confusion_matrix(y_train, y_train_predict_rfc)


# Print Classification Report for training dataset
print('Classification Report for Random Forest (Training Dataset):\n',
      ↪classification_report(y_train, y_train_predict_rfc))
```

Classification Report for Random Forest (Training Dataset):

	precision	recall	f1-score	support
0	1.00	1.00	1.00	180
1	1.00	1.00	1.00	314
2	1.00	1.00	1.00	247
accuracy			1.00	741
macro avg	1.00	1.00	1.00	741
weighted avg	1.00	1.00	1.00	741

```
[52]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot the confusion matrix for the training dataset
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues',
            xticklabels=np.unique(y_train), yticklabels=np.unique(y_train))
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix for Random Forest (Training Dataset)')
plt.show()
```



```
[53]: model = DecisionTreeClassifier()
model.fit(X_train,y_train)
X_test_predict_dtc = model.predict(X_test)
print('The accuracy for Decision Tree Classifier model is (Test Dataset) ',
      metrics.accuracy_score(X_test_predict_dtc,y_test))
model = RandomForestClassifier(n_estimators=100)
model.fit(X_test,y_test)
X_train_predict =model.predict(X_train)
print('The accuracy for Decision Tree Classifier model is (Train Dataset)',
      metrics.accuracy_score(X_train_predict,y_train))
```

```
The accuracy for Decision Tree Classifier model is (Test Dataset)
0.9937106918238994
The accuracy for Decision Tree Classifier model is (Train Dataset)
0.9986504723346828
```

```
[54]: # Import necessary libraries
from sklearn.tree import DecisionTreeClassifier
```

```

from sklearn import tree
import matplotlib.pyplot as plt

# Train a decision tree classifier
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)

# Visualize the decision tree
plt.figure(figsize=(20, 20))
tree.plot_tree(clf, filled=True)
plt.show()

```



[]:

```
[56]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import metrics

# For Decision Tree Classifier
model_dtc = DecisionTreeClassifier()
model_dtc.fit(X_train, y_train)
X_test_predict_dtc = model_dtc.predict(X_test)
print('The accuracy for Decision Tree Classifier model is (Test Dataset): ',
      ↪metrics.accuracy_score(X_test_predict_dtc, y_test))

# Classification report for test dataset
from sklearn.metrics import classification_report
print("Classification Report for Test Dataset:")
print(classification_report(y_test, X_test_predict_dtc))

# For RandomForestClassifier
model_rf = RandomForestClassifier(n_estimators=100)
model_rf.fit(X_train, y_train)
X_train_predict_rf = model_rf.predict(X_train)
print('The accuracy for RandomForestClassifier model is (Train Dataset): ',
      ↪metrics.accuracy_score(X_train_predict_rf, y_train))

# Classification report for train dataset
print("Classification Report for Train Dataset:")
print(classification_report(y_train, X_train_predict_rf))
```

The accuracy for Decision Tree Classifier model is (Test Dataset):

0.9937106918238994

Classification Report for Test Dataset:

	precision	recall	f1-score	support
0	0.99	0.99	0.99	76
1	0.99	0.99	0.99	115
2	1.00	1.00	1.00	127
accuracy			0.99	318
macro avg	0.99	0.99	0.99	318
weighted avg	0.99	0.99	0.99	318

The accuracy for RandomForestClassifier model is (Train Dataset): 1.0

Classification Report for Train Dataset:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	180
1	1.00	1.00	1.00	314
2	1.00	1.00	1.00	247

accuracy			1.00	741
macro avg	1.00	1.00	1.00	741
weighted avg	1.00	1.00	1.00	741

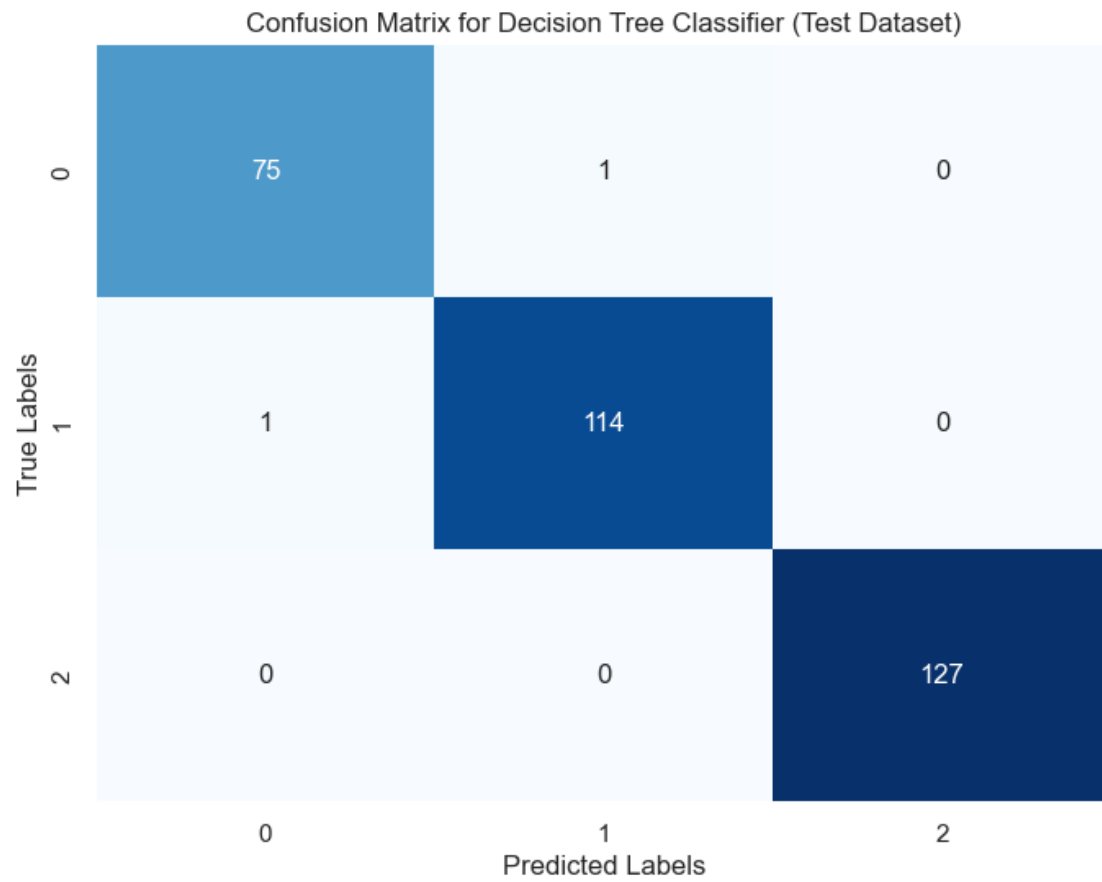
```
[57]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt


# For Decision Tree Classifier
model_dtc = DecisionTreeClassifier()
model_dtc.fit(X_train, y_train)
X_test_predict_dtc = model_dtc.predict(X_test)

# Confusion matrix for test dataset
conf_matrix_test = confusion_matrix(y_test, X_test_predict_dtc)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_test, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=model_dtc.classes_, yticklabels=model_dtc.classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for Decision Tree Classifier (Test Dataset)')
plt.show()

# For RandomForestClassifier
model_rf = RandomForestClassifier(n_estimators=100)
model_rf.fit(X_train, y_train)
X_train_predict_rf = model_rf.predict(X_train)

# Confusion matrix for train dataset
conf_matrix_train = confusion_matrix(y_train, X_train_predict_rf)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=model_rf.classes_, yticklabels=model_rf.classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for RandomForestClassifier (Train Dataset)')
plt.show()
```





[]:

```
[58]: model = KNeighborsClassifier()
model.fit(X_train,y_train)
X_test_predict_knn = model.predict(X_test)
print('The accuracy for KNN classifier model is (Test Dataset)', metrics.
      ↪accuracy_score(X_test_predict_knn,y_test))
model = KNeighborsClassifier()
model.fit(X_test,y_test)
X_train_predict_knn =model.predict(X_train)
print('The accuracy for KNN classifier model is (Train Dataset) ',metrics.
      ↪accuracy_score(X_train_predict_knn,y_train))
```

The accuracy for KNN classifier model is (Test Dataset) 0.9937106918238994
 The accuracy for KNN classifier model is (Train Dataset) 0.9568151147098516

```
[59]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
```

```

# For KNN classifier
model_knn = KNeighborsClassifier()
model_knn.fit(X_train, y_train)
X_test_predict_knn = model_knn.predict(X_test)

# Classification report for test dataset
print("Classification Report for Test Dataset:")
print(classification_report(y_test, X_test_predict_knn))

# For the train dataset
model_knn_train = KNeighborsClassifier()
model_knn_train.fit(X_test, y_test)
X_train_predict_knn = model_knn_train.predict(X_train)

# Classification report for train dataset
print("Classification Report for Train Dataset:")
print(classification_report(y_train, X_train_predict_knn))

```

Classification Report for Test Dataset:

	precision	recall	f1-score	support
0	0.99	1.00	0.99	76
1	1.00	0.98	0.99	115
2	0.99	1.00	1.00	127
accuracy			0.99	318
macro avg	0.99	0.99	0.99	318
weighted avg	0.99	0.99	0.99	318

Classification Report for Train Dataset:

	precision	recall	f1-score	support
0	0.99	0.96	0.98	180
1	0.97	0.92	0.95	314
2	0.91	1.00	0.95	247
accuracy			0.96	741
macro avg	0.96	0.96	0.96	741
weighted avg	0.96	0.96	0.96	741

```

[60]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# For KNN classifier

```

```

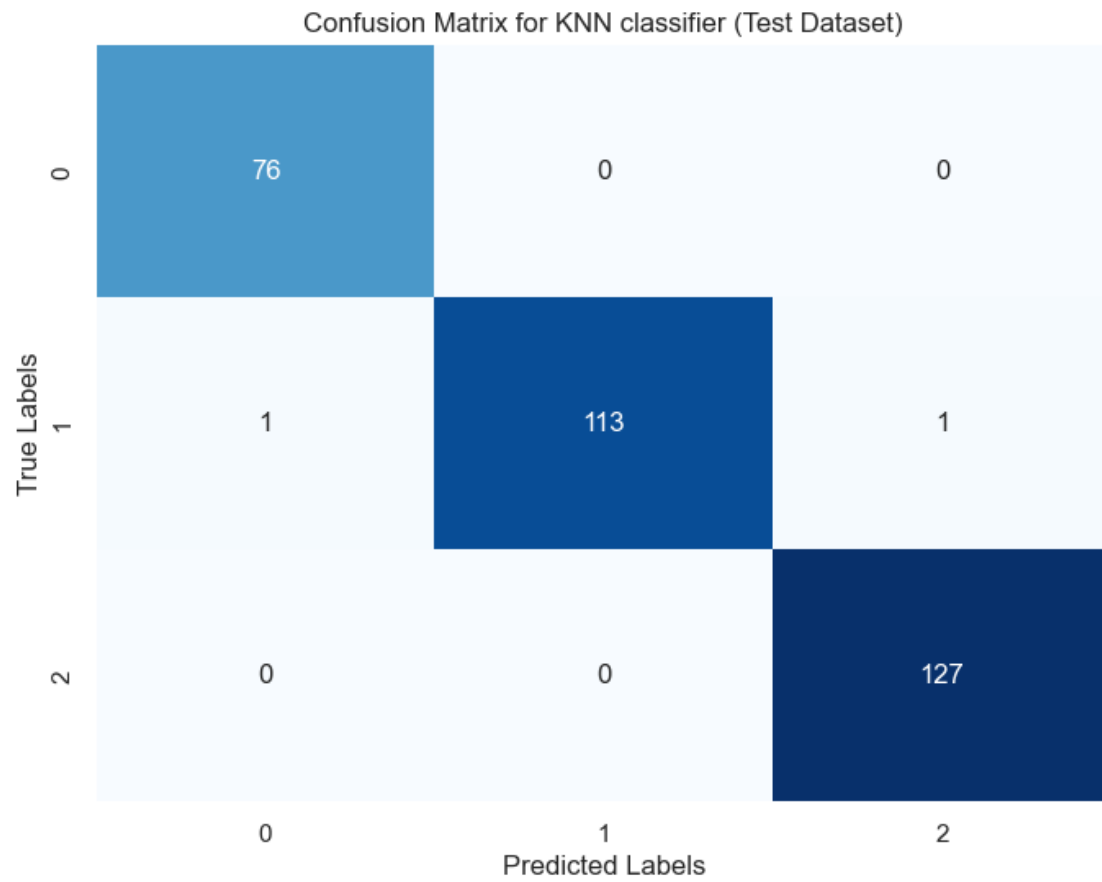
model_knn = KNeighborsClassifier()
model_knn.fit(X_train, y_train)
X_test_predict_knn = model_knn.predict(X_test)


# Confusion matrix for test dataset
conf_matrix_test = confusion_matrix(y_test, X_test_predict_knn)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_test, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=model_knn.classes_, yticklabels=model_knn.classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for KNN classifier (Test Dataset)')
plt.show()

# For the train dataset
model_knn_train = KNeighborsClassifier()
model_knn_train.fit(X_test, y_test)
X_train_predict_knn = model_knn_train.predict(X_train)

# Confusion matrix for train dataset
conf_matrix_train = confusion_matrix(y_train, X_train_predict_knn)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix_train, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=model_knn_train.classes_, yticklabels=model_knn_train.
            ↪classes_)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for KNN classifier (Train Dataset)')
plt.show()

```





```
[61]: test_acc=[]
      train_acc=[]

      for i in range(1,15):
          knn = KNeighborsClassifier(i) #setting up a knn classifier
          knn.fit(X_train,y_train) #fitting the model
          # computing the accuracy for both the trainig and the test data
          train_acc.append(knn.score(X_train,y_train))
          test_acc.append(knn.score(X_test,y_test))
```

```
[62]: import matplotlib.pyplot as plt
      import seaborn as sb

      plt.figure(figsize=(14, 8))
      plt.title('The k-NN Varying number of neighbors')

      # Assuming train_acc and test_acc are lists or arrays containing your data
      sb.lineplot(x=range(1, 15), y=train_acc, label='Training Data Accuracy')
      sb.lineplot(x=range(1, 15), y=test_acc, label='Test Data Accuracy')
```

```
plt.show()
```



```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```